

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [3]: data.describe()
```

```
Out[3]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [4]: data.head(10)
```

```
Out[4]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
5	6	pop	74	3623	70225	1	45.000702	7.682270	7900
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
8	9	sport	73	4049	76000	1	45.548000	11.549470	5600
9	10	sport	51	3653	89000	1	45.438301	10.991700	6000

```
In [5]: data1=data.drop(['lat','ID'],axis=1)
```

```
In [6]: data1
```

```
Out[6]:
```

	model	engine_power	age_in_days	km	previous_owners	lon	price
0	lounge	51	882	25000	1	8.611560	8900
1	pop	51	1186	32500	1	12.241890	8800
2	sport	74	4658	142228	1	11.417840	4200
3	lounge	51	2739	160000	1	17.634609	6000
4	pop	73	3074	106880	1	12.495650	5700
...
1533	sport	51	3712	115280	1	7.704920	5200
1534	lounge	74	3835	112000	1	8.666870	4600
1535	pop	51	2223	60457	1	9.413480	7500
1536	lounge	51	2557	80750	1	7.682270	5990
1537	pop	51	1766	54276	1	17.568270	7900

1538 rows × 7 columns

```
In [7]: data1=pd.get_dummies(data1)
```

```
In [8]: data1.shape
```

```
Out[8]: (1538, 9)
```

```
In [9]: data1
```

```
Out[9]:
```

	engine_power	age_in_days	km	previous_owners	lon	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8.611560	8900	1	0	0
1	51	1186	32500	1	12.241890	8800	0	1	0
2	74	4658	142228	1	11.417840	4200	0	0	1
3	51	2739	160000	1	17.634609	6000	1	0	0
4	73	3074	106880	1	12.495650	5700	0	1	0
...
1533	51	3712	115280	1	7.704920	5200	0	0	1
1534	74	3835	112000	1	8.666870	4600	1	0	0
1535	51	2223	60457	1	9.413480	7500	0	1	0
1536	51	2557	80750	1	7.682270	5990	1	0	0
1537	51	1766	54276	1	17.568270	7900	0	1	0

1538 rows × 9 columns

```
In [10]: y=data1['price']  
x=data1.drop('price',axis=1)
```

In [11]:

y

Out[11]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
```

Name: price, Length: 1538, dtype: int64

In [12]:

```
#!/pip3 install scikit-learn
```

In [13]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
```

In [14]:

```
x_test.head(5)
```

Out[14]:

	engine_power	age_in_days	km	previous_owners	lon	model_lounge	model_pop	model_sport
481	51	3197	120000	2	18.167629	0	1	0
76	62	2101	103000	1	8.644440	0	1	0
1502	51	670	32473	1	14.208810	1	0	0
669	51	913	29000	1	8.946250	1	0	0
1409	51	762	18800	1	9.928310	1	0	0

```
In [15]: y_test.head(10)
```

```
Out[15]: 481      7900
         76      7900
         1502    9400
         669    8500
         1409    9700
         1414    9900
         1089    9900
         1507    9950
         970   10700
         1198    8999
         Name: price, dtype: int64
```

```
In [16]: from sklearn.linear_model import LinearRegression
         reg=LinearRegression()#creating object of linearRegression
         reg.fit(x_train,y_train)#training and fitting LR object using training data
```

```
Out[16]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [17]: ypred=reg.predict(x_test)
```

In [18]: ypred

```
Out[18]: array([ 5744.90385381,  7193.08174957,  9821.47324371,  9773.81447482,
 10071.07869345,  9641.63868511,  9649.51258809, 10157.61850988,
  9847.94985009,  9309.71516629, 10482.47565268,  7779.30441243,
  7673.74056422,  6505.72766141,  9619.50289515, 10411.90973558,
  9670.05131682,  7699.50515859,  4877.54154624, 10452.07538228,
 10421.3266699 , 10430.13589594,  7496.76330979,  9961.06726287,
  7050.37003331,  8984.92007278,  4828.00807556,  6998.7171637 ,
  7878.93692686,  9631.98929627,  7384.85817189,  5253.54744257,
  5442.11536379,  5042.43511877,  8999.34734771,  5716.87620176,
  9811.44986065,  8238.65183735,  6305.39688435,  8391.28310387,
  9759.44418847,  6769.38469788,  9169.54303692, 10154.06276081,
  8635.65508981, 10298.38883653,  9048.46307993,  8857.07716202,
  7059.85863786,  9061.62298677,  9451.88481593, 10320.57884532,
 10145.81082462,  6753.76039054,  9742.26723839,  9421.71535673,
  9510.86366496, 10468.92783011,  9789.73337471,  7192.86742992,
 10015.85984307,  7050.64774752,  9908.93831182,  7147.76311105,
  6439.07590176,  9965.68603975,  9828.44710958,  8554.15295046,
  8504.01172278,  6382.30585686,  7747.58493918,  6835.01812084,
  8314.98295981, 10435.98168268,  7353.94396099,  8551.12772532,
  8806.26450005, 10010.27272600,  7442.47221065,  8457.41251420]
```

In [19]: `from sklearn.metrics import r2_score`
`r2_score(y_test,ypred)`

Out[19]: 0.8426736468631826

In [20]: `from sklearn.metrics import mean_squared_error`
`l=mean_squared_error(ypred,y_test)`

In [21]: l

Out[21]: 577771.1158642182

In [22]: `n=577771.1158642182`
`print(n**(1/2))`

760.1125678899266

In [23]: ypred

```
Out[23]: array([ 5744.90385381,  7193.08174957,  9821.47324371,  9773.81447482,
 10071.07869345,  9641.63868511,  9649.51258809, 10157.61850988,
  9847.94985009,  9309.71516629, 10482.47565268,  7779.30441243,
  7673.74056422,  6505.72766141,  9619.50289515, 10411.90973558,
  9670.05131682,  7699.50515859,  4877.54154624, 10452.07538228,
 10421.3266699 , 10430.13589594,  7496.76330979,  9961.06726287,
  7050.37003331,  8984.92007278,  4828.00807556,  6998.7171637 ,
  7878.93692686,  9631.98929627,  7384.85817189,  5253.54744257,
  5442.11536379,  5042.43511877,  8999.34734771,  5716.87620176,
  9811.44986065,  8238.65183735,  6305.39688435,  8391.28310387,
  9759.44418847,  6769.38469788,  9169.54303692, 10154.06276081,
  8635.65508981, 10298.38883653,  9048.46307993,  8857.07716202,
  7059.85863786,  9061.62298677,  9451.88481593, 10320.57884532,
 10145.81082462,  6753.76039054,  9742.26723839,  9421.71535673,
  9510.86366496, 10468.92783011,  9789.73337471,  7192.86742992,
 10015.85984307,  7050.64774752,  9908.93831182,  7147.76311105,
  6439.07590176,  9965.68603975,  9828.44710958,  8554.15295046,
  8504.01172278,  6382.30585686,  7747.58493918,  6835.01812084,
  8314.98295981, 10435.98168268,  7353.94396099,  8551.12772532,
  8886.26450005, 10010.27272600,  7442.47221065,  8457.41251420,
```



```
In [24]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results.head(15)
```

Out[24]:

	price	predicted
481	7900	5744.903854
76	7900	7193.081750
1502	9400	9821.473244
669	8500	9773.814475
1409	9700	10071.078693
1414	9900	9641.638685
1089	9900	9649.512588
1507	9950	10157.618510
970	10700	9847.949850
1198	8999	9309.715166
1088	9890	10482.475653
576	7990	7779.304412
965	7380	7673.740564
1488	6800	6505.727661
1432	8900	9619.502895

```
In [26]: Results['DIFF']=Results.apply(lambda row:row.price-row.predicted,axis=1)
```

In [27]: Results

Out[27]:

	price	predicted	DIFF
481	7900	5744.903854	2155.096146
76	7900	7193.081750	706.918250
1502	9400	9821.473244	-421.473244
669	8500	9773.814475	-1273.814475
1409	9700	10071.078693	-371.078693
...
291	10900	10039.516366	860.483634
596	5699	6282.912490	-583.912490
1489	9500	9991.170743	-491.170743
1436	6990	8331.810958	-1341.810958
575	10900	10357.579392	542.420608

508 rows × 3 columns

In []: